

Assignment 1: Churn Prediction

May 26, 2024

Group 5

- Moqian Chen (r0965473)
- Laurens Dergent (r0794288)
- Sarah Guilliams (r0751825)
- Yeabsera Kinfu (r0930148)
- Jorge Puertolas Molina (r0978889)
- Isabel Scholz (r1008561)

Link to the GitHub repository: <https://github.com/LaurensDergent/DataAnalytics.git>

0.1 Introduction

In today's competitive business environment, customer retention is critical for the success of any organization. A good churn prediction can help a business sustain customers.

This assignment delves into the prediction of customer churn using diverse machine learning techniques. The analysis at hand is comprised of several steps: the preprocessing of the data to make sure the data is compatible with the models, feature engineering to create new potentially meaningful predictors, and the evaluation of different machine learning models using the custom metric, profit top k. This assignment will also implement a fine-tuning of models to specifically target high scores in the "profit top k" metric. This involves not only optimizing model parameters through hyperparameter tuning, but also using sample weights to enhance the model's ability to identify high-spending customers who are most likely to churn.

```
[1]: import pandas as pd
from datetime import datetime
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from collections import defaultdict
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import MinMaxScaler
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
import matplotlib.pyplot as plt
```

```

import numpy as np
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from sklearn.preprocessing import LabelEncoder
#import featuretools as ft
from sklearn.model_selection import KFold
from sklearn.model_selection import StratifiedKFold, GridSearchCV
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import StackingRegressor
import os
from sklearn.inspection import PartialDependenceDisplay
import warnings

from sklearn.model_selection import cross_val_score

from sklearn.metrics import confusion_matrix, roc_curve, auc, roc_auc_score
from sklearn.linear_model import LogisticRegression

# To ignore all warnings
warnings.filterwarnings("ignore")

# To ignore specific warnings by category
# For example, to ignore all RuntimeWarning
warnings.filterwarnings("ignore", category=RuntimeWarning)

trainval = pd.read_csv("train.csv")
test = pd.read_csv("test.csv")

```

Before proceeding further, a dictionary was created to personalize the pipeline and experiment with various combinations of methods for model training. This dictionary contains parameters that can be adjusted for preprocessing purposes. These parameters include featuretools, undersampling, oversampling, timecolumn, corrthresh, and corrthresh2.

The featuretools parameter decides whether the Featuretools package is used for feature engineering or not. The timecolumn parameter specifies whether the column of connection time of the user is included in the feature set. Undersampling and oversampling parameters control whether these techniques will be employed to address the class imbalance. The corrthresh and corrthresh2 parameters represent correlation thresholds for dropping variables. All these parameters will be further explained later in the analysis.

```
[2]: extra_params = {"feature_tools" : 1, "undersampling" : 1, "oversampling" : 1,
    ↪ "time_column" : 1, "corr_thresh" : 0.9, "corr_thresh2" : 0.9}
```

0.2 Division of Sets

The initial step in the pipeline was partitioning the datasets. This entails separating the features from the dependent variable, as well as splitting the original training data into two subsets: the training set and the validation set. In this case, the validation set would be the equivalent of a test set, since the test set does not contain labels. This division is crucial for evaluating the models. All models were tested on the validation set, and then the best-performing one was selected for the leaderboard. The validation set was chosen to be 20 % of the original training set.

```
[2]: X_trainval = trainval.drop(columns = "target")
    y_trainval = trainval["target"]
```

```
[3]: train, val, train_target, val_target = train_test_split(X_trainval, y_trainval,
    ↪ test_size=0.2, random_state=42)
```

```
[4]: train['target'] = train_target
    val['target'] = val_target
```

It was necessary to store the minimum average cost values separately because they are required for computing the top k metric. When scaling features, the original values of average cost min were lost, and this was a way of gathering them again when computing the metric.

```
[5]: #Keep average cost min
    averagecostmin = val['average cost min']
    idtest = test["id"]
```

0.3 Initial Exploration

To start, we visualize the distributions of numerical variables. This is important to see whether any transformations should be applied to these variables. First, we can notice that the scaling of the variables vary a lot. PeakminsSum's maximum value is 2901, while WeekendRatio goes up to 0.74. This is important to take into account when inputting the data in models that are vulnerable to changes in the scaling of the variables. Second of all, most variables are not symmetric, implying they are not normally distributed. Most 75th quantiles are much farther from the median than 25th quantiles are. In this prediction environment where extreme values are likely to cause churnings, normalizing these variables was not considered to be suitable. High values could be the thing that causes churning and seeing this drastic difference between observations could help the model discriminate. By normalizing we make the line between high and medium values finer.

```
[6]: train.describe()
```

```
[6]:
```

	Age	L_0_S	Dropped_Calls	Peak_calls_Sum	Peak_mins_Sum	\
count	4035.000000	4035.000000	4035.000000	4035.000000	4035.000000	
mean	31.414126	33.699347	2.673358	238.945229	708.792184	
std	12.775443	14.007828	3.418703	239.288923	503.929355	
min	12.000000	9.633333	0.000000	0.000000	0.000000	

25%	22.000000	21.300000	0.000000	59.000000	307.500000
50%	29.000000	33.666667	1.000000	161.000000	614.400001
75%	39.000000	45.933333	2.000000	343.000000	1014.300000
max	80.000000	58.200000	15.000000	1626.000000	2901.600001

	OffPeak_calls_Sum	OffPeak_mins_Sum	Weekend_calls_Sum	\
count	4035.000000	4035.000000	4035.000000	
mean	104.036431	311.938372	16.322677	
std	97.260823	199.167646	16.296733	
min	0.000000	0.000000	0.000000	
25%	28.000000	149.550000	4.000000	
50%	73.000000	291.000000	11.000000	
75%	156.000000	446.400000	24.000000	
max	560.000000	1091.099999	106.000000	

	Weekend_mins_Sum	International_mins_Sum	...	call_cost_per_min	\
count	4035.000000	4035.000000	...	4032.000000	
mean	50.079971	169.551549	...	10.081397	
std	36.092455	141.787096	...	2.108023	
min	0.000000	0.000000	...	2.000000	
25%	22.800000	64.261044	...	8.573250	
50%	44.399999	132.279506	...	9.518597	
75%	71.750000	236.853610	...	11.696253	
max	205.000000	935.947864	...	21.734694	

	actual call cost	Total_call_cost	Total_Cost	average cost min	\
count	4035.000000	4035.000000	4035.000000	4035.000000	
mean	19.133955	73.501256	186.803427	0.168306	
std	27.987193	60.000356	82.079489	0.077812	
min	0.000000	0.000000	59.940000	0.048998	
25%	0.000000	32.134081	123.802373	0.134413	
50%	7.515854	57.801016	173.254831	0.154525	
75%	28.308661	99.217719	232.116107	0.179302	
max	184.892166	437.063835	587.063835	1.357564	

	Peak ratio	OffPeak ratio	Weekend ratio	Nat-InterNat Ratio	\
count	4035.000000	4035.000000	4035.000000	4035.000000	
mean	0.609716	0.327824	0.061717	0.162328	
std	0.224430	0.209932	0.063644	0.105643	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.466317	0.165139	0.020942	0.078918	
50%	0.654877	0.287722	0.043932	0.157734	
75%	0.782261	0.458395	0.080924	0.244154	
max	1.000000	1.000000	0.731884	2.601223	

	target
count	4035.000000

```

mean      0.146716
std       0.353867
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       1.000000

```

[8 rows x 30 columns]

```
[7]: train.describe(include = 'object')
```

```

[7]:      Gender  Connect_Date  tariff  Handset  Usage_Band  Tariff_OK  \
count      4035           4035      4035      4035          4032      4035
unique         2           1310         5         11           5         4
top           F      11/07/99  CAT 200       S50         Med         OK
freq       2042             11      1802       944       2232       4006

      high Dropped calls No Usage      id
count           4035      4035      4035
unique            2         2      4035
top              F         F  K277140
freq           3933      4032         1

```

The table above was used to delete redundant variables. Along this line, the id variable was deleted. "Usage" was deleted from the model because it contributed little to no information, since out of 4035 observations, it had 4032 repetitions of "F". "Tariff OK" also showed the same characteristics, and had 4006 repetitions of "OK". However, it was retained because the remaining observations could still capture information on whether the client would churn or not.

```

[8]: train = train.drop(columns = ["id"])
      val = val.drop(columns = ["id"])
      test = test.drop(columns = ["id"])

```

0.4 Preprocessing

Feature Engineering A new variable called "ConnectTime" was created, that represented the amount of time the user has been connected for/affiliated with the company. This was hypothesized to be useful as it was a proxy of loyalty. For generating this variable, the values in the ConnectDate column were subtracted from the last date in the train set. Using the last date in the train set ensured that the cutoff point was the same for all the instances across all sets.

```

[9]: if extra_params["time_column"] == 1:
      train["Connect_Date"] = train["Connect_Date"].apply(lambda x : datetime.
      ↳strptime(x, "%d/%m/%y"))
      test["Connect_Date"] = test["Connect_Date"].apply(lambda x : datetime.
      ↳strptime(x, "%d/%m/%y"))

```

```

val["Connect_Date"] = val["Connect_Date"].apply(lambda x : datetime.
↳strptime(x, "%d/%m/%y"))

max_train = train["Connect_Date"].max()

train["End_Date"] = max_train
test["End_Date"] = max_train
val["End_Date"] = max_train

train["Connect_Time"] = train["End_Date"] - train["Connect_Date"]
test["Connect_Time"] = test["End_Date"] - test["Connect_Date"]
val["Connect_Time"] = val["End_Date"] - val["Connect_Date"]

train["Connect_Time"] = train["Connect_Time"].apply(lambda x : x.days)
test["Connect_Time"] = test["Connect_Time"].apply(lambda x : x.days)
val["Connect_Time"] = val["Connect_Time"].apply(lambda x : x.days)

train = train.drop(columns = ["End_Date"])
test = test.drop(columns = ["End_Date"])
val = val.drop(columns = ["End_Date"])

train = train.drop(columns = ["Connect_Date"])
test = test.drop(columns = ["Connect_Date"])
val = val.drop(columns = ["Connect_Date"])

```

Missing Values

```
[10]: train.isna().sum()
```

```

[10]: Gender                0
      Age                  0
      L_O_S                0
      Dropped_Calls        0
      tariff               0
      Handset              0
      Peak_calls_Sum       0
      Peak_mins_Sum        0
      OffPeak_calls_Sum    0
      OffPeak_mins_Sum     0
      Weekend_calls_Sum    0
      Weekend_mins_Sum     0
      International_mins_Sum 0
      Nat_call_cost_Sum    0
      AvePeak              0
      AveOffPeak           0
      AveWeekend           0
      National_calls       0
      National mins        0

```

```

AveNational          0
All_calls_mins       0
Dropped_calls_ratio  3
Usage_Band           3
Mins_charge          0
call_cost_per_min    3
actual call cost     0
Total_call_cost      0
Total_Cost           0
Tariff_OK            0
average cost min     0
Peak ratio           0
OffPeak ratio        0
Weekend ratio        0
Nat-InterNat Ratio   0
high Dropped calls   0
No Usage             0
target              0
Connect_Time         0
dtype: int64

```

```
[11]: rows_with_missing_values = train[train.isnull().any(axis=1)]
```

```
[12]: rows_with_missing_values[["call_cost_per_min", "Dropped_calls_ratio",
    ↪ "Usage_Band"]]
```

```
[12]:      call_cost_per_min  Dropped_calls_ratio  Usage_Band
3836                NaN                NaN        NaN
4301                NaN                NaN        NaN
3237                NaN                NaN        NaN
```

Out of 4035 observations, only three rows contain missing values. Therefore, a simple method sufficed for imputation. Median and most-frequent imputation methods were implemented using the SimpleImputer package from sklearn.

Imputing Numerical Columns

```
[13]: # Assuming 'data' is your dataset with missing values
# Create an instance of SimpleImputer with strategy='median'
imputer = SimpleImputer(strategy='median')

# Fit the imputer on the data
imputer.fit(train[["call_cost_per_min", "Dropped_calls_ratio"]])

# Transform the data by replacing missing values with the median
test_clean = imputer.transform(test[["call_cost_per_min",
    ↪ "Dropped_calls_ratio"]])
```

```
train_clean = imputer.transform(train[["call_cost_per_min",
↪ "Dropped_calls_ratio"]])
val_clean = imputer.transform(val[["call_cost_per_min", "Dropped_calls_ratio"]])
```

```
[14]: train["call_cost_per_min"] = train_clean[:, 0]
test["call_cost_per_min"] = test_clean[:, 0]
val["call_cost_per_min"] = val_clean[:, 0]

train["Dropped_calls_ratio"] = train_clean[:, 1]
test["Dropped_calls_ratio"] = test_clean[:, 1]
val["Dropped_calls_ratio"] = val_clean[:, 1]
```

Imputing Categorical Column

```
[15]: # Initialize the imputer with 'most_frequent' strategy
imputer = SimpleImputer(strategy='most_frequent')

imputer.fit(np.array(train["Usage_Band"]).reshape(-1, 1))

train["Usage_Band"] = imputer.transform(np.array(train["Usage_Band"]).
↪ reshape(-1, 1))[:, 0]
test["Usage_Band"] = imputer.transform(np.array(test["Usage_Band"]).reshape(-1,
↪ 1))[:, 0]
val["Usage_Band"] = imputer.transform(np.array(val["Usage_Band"]).reshape(-1,
↪ 1))[:, 0]
```

Variable Selection

Dropping numerical variables In the process of refining the features of the training set, an analysis was conducted to examine the correlation among numerical variables. A correlation threshold (corrthresh) of 0.9 was employed to identify pairs of variables with significant correlation. Subsequently, these correlated variable pairs were grouped into clusters such that all connected pairs were in the same cluster. Then, within each cluster, the variable with the highest correlation with the target variable was selected for the final feature set. Hence, the higher the chosen correlation threshold, the less the number of variables that are dropped. The chosen correlation threshold (corrthresh) was high to ensure that valuable variables were not discarded.

```
[16]: correlation_matrix = train[train.select_dtypes(include=['number']).columns.
↪ tolist()].corr()
```

```
[17]: corrtarget = correlation_matrix['target']
```

```
[18]: # Make plot for this
```

```
[19]: high_correlation_pairs = []
for i in range(len(correlation_matrix.columns)):
    for j in range(i+1, len(correlation_matrix.columns)):
        if abs(correlation_matrix.iloc[i, j]) > extra_params["corr_thresh"]:
```



```
        high_correlation_pairs.append((correlation_matrix.columns[i],  
↪correlation_matrix.columns[j]))
```

```
[20]: high_correlation_pairs
```

```
[20]: [('L_O_S', 'Connect_Time'),  
      ('Peak_calls_Sum', 'National_calls'),  
      ('Peak_mins_Sum', 'National mins'),  
      ('Peak_mins_Sum', 'All_calls_mins'),  
      ('International_mins_Sum', 'Total_call_cost'),  
      ('Nat_call_cost_Sum', 'actual call cost'),  
      ('National mins', 'All_calls_mins'),  
      ('All_calls_mins', 'Total_Cost'),  
      ('Total_call_cost', 'Total_Cost'),  
      ('Peak ratio', 'OffPeak ratio')]
```

```
[21]: def build_adjacency_list(pairs):  
    adjacency_list = defaultdict(list)  
    for u, v in pairs:  
        adjacency_list[u].append(v)  
        adjacency_list[v].append(u)  
    return adjacency_list  
  
    # Function to perform depth-first search (DFS) traversal to find connected_  
    ↪components  
    def dfs(node, adjacency_list, visited, component):  
        visited.add(node)  
        component.append(node)  
        for neighbor in adjacency_list[node]:  
            if neighbor not in visited:  
                dfs(neighbor, adjacency_list, visited, component)  
  
    # Function to find connected components in the graph  
    def find_connected_components(pairs):  
        adjacency_list = build_adjacency_list(pairs)  
        visited = set()  
        connected_components = []  
        for node in adjacency_list:  
            if node not in visited:  
                component = []  
                dfs(node, adjacency_list, visited, component)  
                connected_components.append(component)  
        return connected_components  
  
    # Find connected groups  
    connected_groups = find_connected_components(high_correlation_pairs)
```

```
[22]: def get_variables_to_drop(group, corrtarget):
    max_cor = 0
    var_max = ''
    for var in group:
        if abs(corrtarget[var]) > max_cor:
            var_max = var
            max_cor = abs(corrtarget[var])
    group.remove(var_max)

    return group
```

```
[23]: variables_to_drop = [get_variables_to_drop(group, corrtarget) for group in
    ↪connected_groups]
```

```
[24]: drop_vars = [item for sublist in variables_to_drop for item in sublist]
```

```
[25]: train = train.drop(columns = drop_vars)
```

```
[26]: val = val.drop(columns = drop_vars)
```

```
[27]: test = test.drop(columns = drop_vars)
```

Dropping categorical variables

```
[28]: for variable in list(train.describe(include = 'object').columns):
    print(train.groupby('target')[variable].value_counts())
```

```
target  Gender
0        F      1761
         M      1682
1        M       311
         F       281
Name: Gender, dtype: int64
target  tariff
0      CAT 200    1533
         CAT 100     736
         Play 100    503
         Play 300    471
         CAT 50     200
1      CAT 200     269
         Play 100    128
         CAT 100     110
         Play 300     47
         CAT 50      38
Name: tariff, dtype: int64
target  Handset
0      S50       825
         BS110    599
         S80      574
```

	WC95	526
	ASAD170	514
	BS210	220
	CAS60	90
	ASAD90	42
	CAS30	37
	SOP20	10
	SOP10	6
1	ASAD90	174
	S50	119
	CAS30	108
	BS110	95
	SOP10	28
	SOP20	24
	S80	20
	ASAD170	10
	BS210	8
	WC95	6

Name: Handset, dtype: int64

	target	Usage_Band
0	Med	1974
	MedHigh	866
	MedLow	310
	High	246
	Low	47
1	Med	261
	MedHigh	148
	MedLow	120
	High	42
	Low	21

Name: Usage_Band, dtype: int64

	target	Tariff_OK
0	OK	3436
	High CAT 100	5
	High CAT 50	1
	High Play 100	1
1	OK	570
	High CAT 100	16
	High CAT 50	4
	High Play 100	2

Name: Tariff_OK, dtype: int64

	target	high Dropped calls
0	F	3426
	T	17
1	F	507
	T	85

Name: high Dropped calls, dtype: int64

target No Usage

0	F	3442
	T	1
1	F	590
	T	2

Name: No Usage, dtype: int64

The variables UsageBand, tariff, and NoUsage show no valuable information to predict the target variable. Their distributions do not differentiate between churners and non-churners, suggesting minimal discriminatory power. However, High Dropped Calls seems like a good indicator for predicting churn, with a notably higher ratio of True values among churners compared to non-churners.

Furthermore, the variable TariffOk, particularly in the category of 100, shows a substantially higher presence among churners compared to non-churners. Encoding this variable using One-Hot encoding could enhance its predictive value.

Similarly, considering the variability in handset types across churning and non-churning groups, using one-hot encoding for handset appears to be a suitable option. This approach would enable the model to capture the nuances associated with different handsets and potentially improve predictive accuracy.

```
[29]: train = train.drop(columns = ["No Usage", "Usage_Band", "tariff"])
      test = test.drop(columns = ["No Usage", "Usage_Band", "tariff"])
      val = val.drop(columns = ["No Usage", "Usage_Band", "tariff"])
```

Outliers are not dealt with since they are useful in the detection of churners. Churning could be the precise thing that is caused by outliers.

Variable Transformation

Numerical Variables To make the numerical variables easier for the models, the numerical variables were transformed to a 0-1 scale using MinMaxScaler. As presented in the Initial Exploration, normalized was rejected since churning prediction benefits from extremes.

```
[31]: # List of numerical variable names
numerical_features = train.select_dtypes(include=['number']).columns.tolist()

# List of categorical variable names (replace with actual categorical variable_
↳ names)
categorical_features = train.select_dtypes(exclude=['number']).columns.tolist()

# Define the transformers
transformers = [
    ('num', MinMaxScaler(), numerical_features)
]

# Create the column transformer
preprocessor = ColumnTransformer(transformers, remainder='passthrough')

# Apply the column transformer to the data
```

```

transformed_train = preprocessor.fit_transform(train)
transformed_val = preprocessor.transform(val)
transformed_test = preprocessor.transform(test)

# Convert the transformed data back to DataFrame
transformed_train = pd.DataFrame(transformed_train, columns=numerical_features +
    ↪categorical_features)
transformed_test = pd.DataFrame(transformed_test, columns=numerical_features +
    ↪categorical_features)
transformed_val = pd.DataFrame(transformed_val, columns=numerical_features +
    ↪categorical_features)

# Only the numerical variables will be scaled to the range [0, 1]

```

Categorical Variables Following the insights from the Variable Selection section, categorical variables were transformed using One-Hot Encoding, although not all levels of each categorical variable were included. Specifically, for TariffOk, only a binary variable indicating whether it was HighCat100 or not was added.

```

[32]: for variable in list(train.describe(include = 'object').columns):
        print(train.groupby('target')[variable].value_counts())

```

```

target  Gender
0       F      1761
       M      1682
1       M       311
       F       281
Name: Gender, dtype: int64
target  tariff
0       CAT 200   1533
       CAT 100    736
       Play 100    503
       Play 300    471
       CAT 50     200
1       CAT 200    269
       Play 100    128
       CAT 100    110
       Play 300     47
       CAT 50     38
Name: tariff, dtype: int64
target  Handset
0       S50      825
       BS110     599
       S80      574
       WC95     526
       ASAD170  514
       BS210    220

```

	CAS60	90
	ASAD90	42
	CAS30	37
	SOP20	10
	SOP10	6
1	ASAD90	174
	S50	119
	CAS30	108
	BS110	95
	SOP10	28
	SOP20	24
	S80	20
	ASAD170	10
	BS210	8
	WC95	6

Name: Handset, dtype: int64

target	Usage_Band	
0	Med	1974
	MedHigh	866
	MedLow	310
	High	246
	Low	47
1	Med	261
	MedHigh	148
	MedLow	120
	High	42
	Low	21

Name: Usage_Band, dtype: int64

target	Tariff_OK	
0	OK	3436
	High CAT 100	5
	High CAT 50	1
	High Play 100	1
1	OK	570
	High CAT 100	16
	High CAT 50	4
	High Play 100	2

Name: Tariff_OK, dtype: int64

target	high Dropped calls	
0	F	3426
	T	17
1	F	507
	T	85

Name: high Dropped calls, dtype: int64

target	No Usage	
0	F	3442
	T	1
1	F	590

T 2
Name: No Usage, dtype: int64

```
[33]: # Encode binary variables using map
binary_mapping1 = {'M': 1, 'F': 0}
binary_mapping2 = {'T': 1, 'F': 0}

transformed_train['high Dropped calls'] = transformed_train['high Dropped_
↳calls'].map(binary_mapping2)
transformed_train['No Usage'] = transformed_train['No Usage'].
↳map(binary_mapping2)
transformed_train['Gender'] = transformed_train['Gender'].map(binary_mapping1)
transformed_test['high Dropped calls'] = transformed_test['high Dropped calls'].
↳map(binary_mapping2)
transformed_test['No Usage'] = transformed_test['No Usage'].map(binary_mapping2)
transformed_test['Gender'] = transformed_test['Gender'].map(binary_mapping1)
transformed_val['high Dropped calls'] = transformed_val['high Dropped calls'].
↳map(binary_mapping2)
transformed_val['No Usage'] = transformed_val['No Usage'].map(binary_mapping2)
transformed_val['Gender'] = transformed_val['Gender'].map(binary_mapping1)

# Filter and encode categorical variables
# For categorical_1, keep only 'High CAT 100', encode the rest as 'Other'
transformed_train['Tariff_OK'] = transformed_train['Tariff_OK'].apply(lambda x:
↳1 if x == 'High CAT 100' else 0)
transformed_test['Tariff_OK'] = transformed_test['Tariff_OK'].apply(lambda x: 1
↳if x == 'High CAT 100' else 0)
transformed_val['Tariff_OK'] = transformed_val['Tariff_OK'].apply(lambda x: 1 if
↳x == 'High CAT 100' else 0)

# Encode categorical_2 using one-hot encoding
encoded_train = pd.get_dummies(transformed_train, columns=['Handset', "tariff",
↳"Usage_Band"], dtype = int)
encoded_test = pd.get_dummies(transformed_test, columns=['Handset', "tariff",
↳"Usage_Band"], dtype = int)
encoded_val = pd.get_dummies(transformed_val, columns=['Handset', "tariff",
↳"Usage_Band"], dtype = int)

[34]: encoded_val['target'] = encoded_val['target'].astype('category')

[35]: X_train = encoded_train.drop(columns = "target")
y_train = encoded_train["target"]

[36]: X_val = encoded_val.drop(columns = "target")
y_val = encoded_val["target"]
```

```
[37]: X_test = encoded_test.drop(columns = "target")
```

```
[38]: # Initialize LabelEncoder
label_encoder = LabelEncoder()

# Fit and transform the categorical labels
y_train = label_encoder.fit_transform(y_train)
y_val = label_encoder.transform(y_val)
```

Oversampling / Undersampling There was a significant class imbalance, the number of churners was much lower than the number of non-churners. To fix this, oversampling and undersampling were implemented.

Oversampling

```
[39]: if extra_params["oversampling"] == 1:
    if X_train is not None and y_train is not None:
        # Ensure y_train is an array of integers
        #print("y_train data type:", type(y_train))

        # Apply SMOTE for oversampling
        smote = SMOTE(random_state=42)
        X_resampled, y_resampled = smote.fit_resample(X_train, y_train)

        # Update X_train and y_train with resampled data
        X_train = X_resampled
        y_train = y_resampled
    else:
        print("X_train or y_train is None. Please check the input data.")
else:
    print("Oversampling is disabled.")
```

Undersampling

```
[40]: # Apply RandomUnderSampler for undersampling
if extra_params["undersampling"] == 1:
    rus = RandomUnderSampler(random_state=42)
    X_resampled, y_resampled = rus.fit_resample(X_train, y_train)

    X_train = X_resampled
    y_train = y_resampled
else:
    print("Undersampling is disabled.")
```

Feature tools In the final preprocessing stage, feature engineering was performed using feature-tools. Multiplication primitives were applied to generate interaction effects among features. Initially, the resulting matrix had approximately 2000 features.

Subsequently, a filtering process was implemented, starting with the elimination of variables containing only one unique value. Following this, variables uncorrelated with the target variable were dropped. Finally, another correlation threshold was applied to further cut down features, leading to a set of 82 variables.

```
[41]: from featuretools.selection import remove_highly_correlated_features
import featuretools as ft

if extra_params['feature_tools'] == 1:
    combined_df = pd.concat([X_train, X_val, X_test], ignore_index=True)
    averagecostmin = combined_df['average cost min']
    es = ft.EntitySet(id='EntitySet')

    # Add your DataFrame as an entity to the EntitySet
    es = es.add_dataframe(
        dataframe_name="AllDataFT",
        dataframe=combined_df,
        index = "index"
    )

    # Perform Deep Feature Synthesis (DFS)
    feature_matrix, feature_defs = ft.dfs(entityset=es,
    ↪target_dataframe_name='AllDataFT',
                                     trans_primitives=['multiply_numeric'],
                                     max_depth=1)

    # Split the combined dataset back into train, validation, and test sets
    train_size = len(X_train)
    val_size = len(X_val)

    X_train = feature_matrix[:train_size]
    X_val = feature_matrix[train_size:train_size + val_size]
    X_test = feature_matrix[train_size + val_size:]
    X_train['target'] = y_train
    #Delete all variables with one unique value
    drop_vars_zero = []
    for var in X_train.columns:
        if len(X_train[var].unique()) == 1:
            drop_vars_zero.append(var)
    X_train = X_train.drop(columns = drop_vars_zero)
    X_val = X_val.drop(columns = drop_vars_zero)
    X_test = X_test.drop(columns = drop_vars_zero)
    #Drop variables (there are too many >2000)
    #Start by dropping uncorrelated variables to target
    correlation_matrix_init = X_train[X_train.select_dtypes(include=['number'])].
    ↪columns.tolist()).corr()
    corrtarget_init = correlation_matrix_init['target'].reset_index()
```

```

drop_init = []
for i in range(0, len(corrtarget_init)):
    if abs(corrtarget_init['target'][i]) <= 0.05:
        drop_init.append(corrtarget_init['index'][i])
X_train = X_train.drop(columns = drop_init)
X_val = X_val.drop(columns = drop_init)
X_test = X_test.drop(columns = drop_init)

threshold = extra_params[corrthresh2]
# Remove highly correlated features
X_train = remove_highly_correlated_features(X_train, pct_corr_threshold = threshold)
X_val = X_val[X_train.columns]
X_test = X_test[X_train.columns]
#X_train['average cost min'] = averagecostmin[:train_size]
#X_val['average cost min'] = averagecostmin[train_size:train_size + val_size]
#X_test['average cost min'] = averagecostmin[train_size+val_size:]

```

```

[42]: X_val = X_val.reset_index(drop = True)
      combined_df['average cost min']

```

```

[42]: 0      0.040380
      1      0.091945
      2      0.061052
      3      0.052613
      4      0.069484
      ...
      9572    0.056049
      9573    0.087503
      9574    0.059979
      9575    0.119850
      9576    0.046733
      Name: average cost min, Length: 9577, dtype: float64

```

0.5 top_20 metric

Once an estimator is trained, the function would directly output the aggregation of top 20 profits from the true positive cases.

```

[43]: def profit_top_20(X_val, y_pred, y_val): ##### PLEASE LET'S KEEP IT THIS WAY
      FOR THE REMAINING TIME!!

      true_positives = []
      for idx, (pred, true_label) in enumerate(zip(y_pred, y_val)):
          if pred == 1 and true_label == 1:
              true_positives.append(idx)

```

```

cost_sort = pd.DataFrame()
cost = X_val['average cost min'][true_positives]

cost_sort['cost']=cost

cost_sort.sort_values('cost', ascending=False, inplace=True)
top_20_sum = cost_sort.head(20).sum()

return top_20_sum

```

0.6 Assign weights to samples

Would-be churners with higher cost are of more profitability. Hence, weights are assigned to samples as clients of higher values are more focal. Weights are determined in an exponential way instead of a linear way, indicating that observations with higher cost would get way more attention in the models.

```

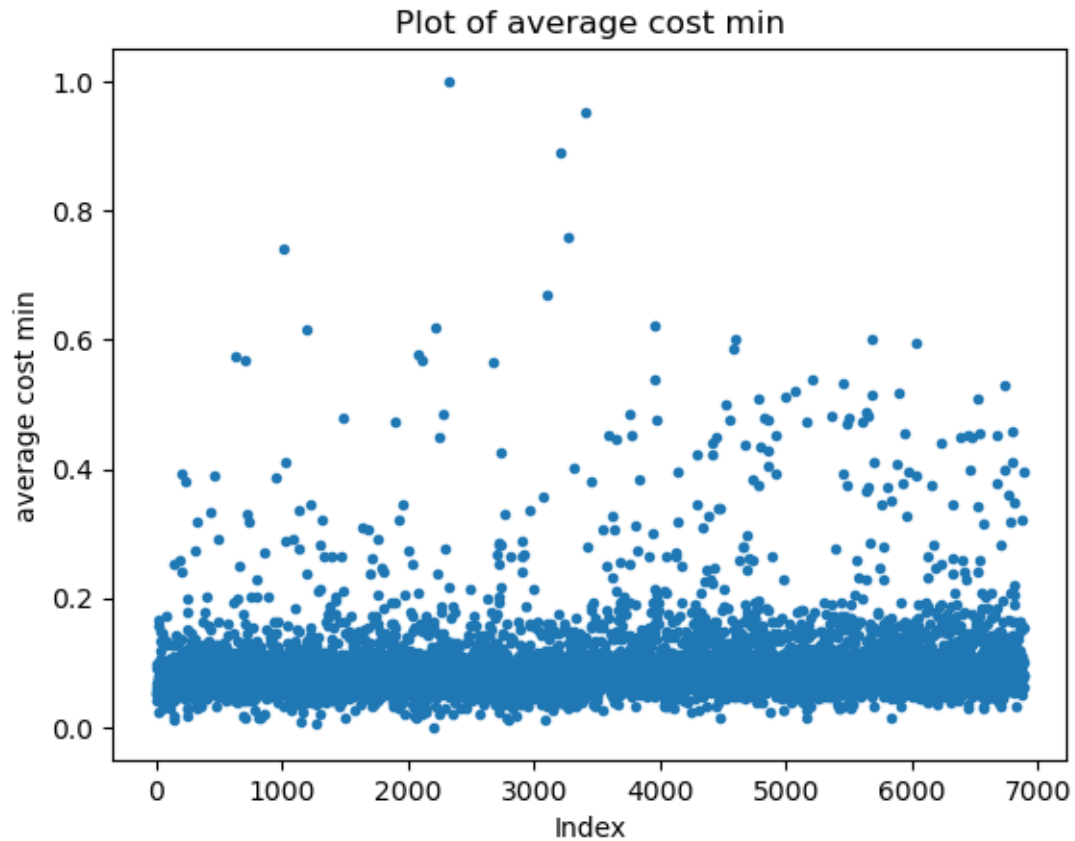
[44]: X_val1,y_val1 = val, val_target
y_train_df = pd.DataFrame(y_train, columns=['target'])
Combined_set = pd.concat([X_train, y_train_df], axis = 1 )
Combined_set.sort_values(by='average cost min')
Combined_set_sorted = Combined_set.sort_values(by='average cost min', ↵
↵ascending=False)

```

```

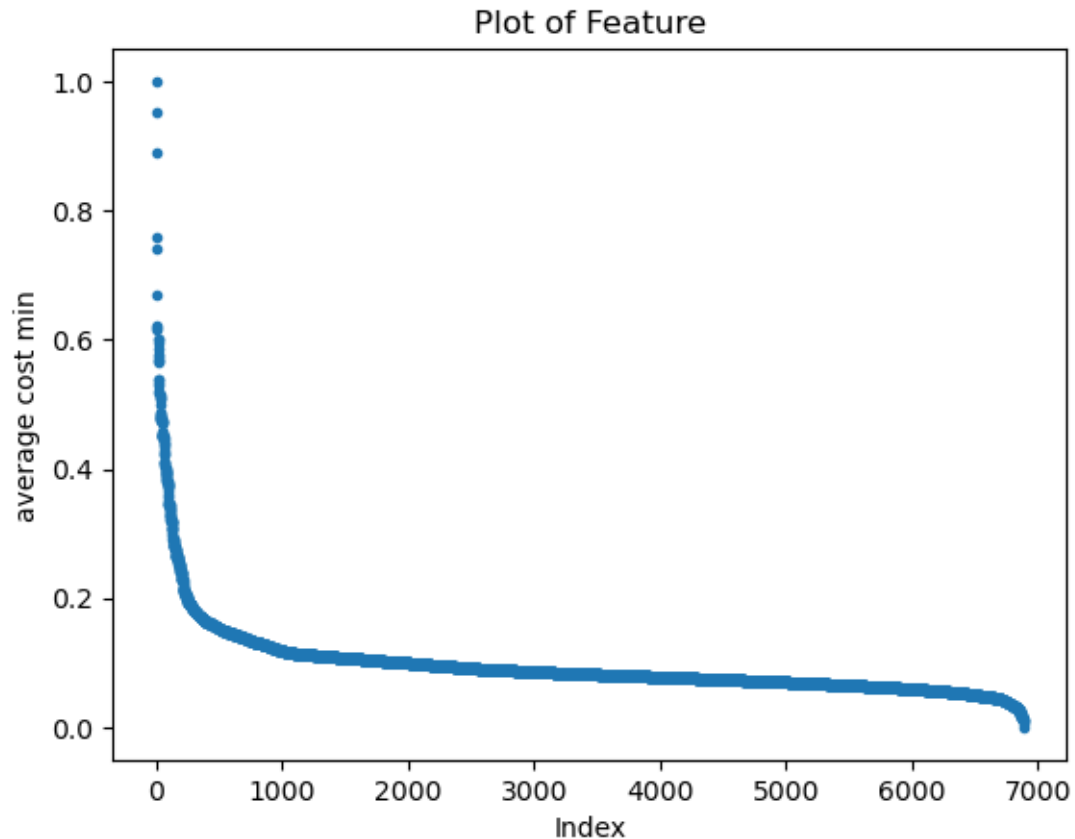
[86]: X_val1,y_val1 = val, val_target
y_train_df = pd.DataFrame(y_train, columns=['target'])
Combined_set = pd.concat([X_train, y_train_df], axis = 1 )
Combined_set.sort_values(by='average cost min')
Combined_set_sorted = Combined_set.sort_values(by='average cost min', ↵
↵ascending=False)
plt.scatter(Combined_set.index, Combined_set["average cost min"], marker='.')
plt.ylabel("average cost min")
plt.xlabel("Index")
plt.title("Plot of average cost min")
plt.show()

```



```
[47]: Combined_set_sorted = Combined_set_sorted.reset_index(drop=True)
X_train_filt = Combined_set_sorted.drop(columns = "target")
y_train_filt = Combined_set_sorted["target"]

[48]: plt.scatter(X_train_filt.index, X_train_filt["average cost min"], marker='.')
plt.ylabel("average cost min")
plt.xlabel("Index")
plt.title("Plot of Feature")
plt.show()
```



```
[85]: # Filter Combined_set_sorted to include only rows where 'target' is equal to 1
filtered_data = Combined_set_sorted[Combined_set_sorted["target"] == 1]

# Select the first 20 rows
filtered_data_first_20 = filtered_data.head(20)

# Get the target values from the filtered DataFrame
y_train_filt20 = filtered_data_first_20["target"]
#print(y_train_filt20)
def profit_top_20_scrap(X_train_filt,pred, y_val):

    true_positives = []
    for idx, (pred, true_label) in enumerate(zip(pred, y_val)):
        if pred == 1 and true_label == 1:
            true_positives.append(idx)

    cost_sort = pd.DataFrame()
    cost = X_train_filt['average cost min'][true_positives]

    cost_sort['cost']=cost
```

```

cost_sort.sort_values('cost', ascending=False, inplace=True)
#print(cost_sort)
top_20_sum = cost_sort.head(20).sum()

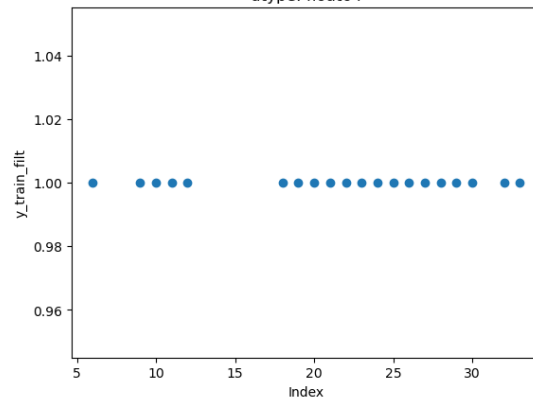
return top_20_sum

##If model perfect our topk is
perf= profit_top_20_scrap(X_train_filt, y_train_filt20, y_train_filt20)

# Plot a scatter plot of y_train_filt against the index
plt.scatter(filtered_data_first_20.index, y_train_filt20)
plt.title(f'Scatter Plot of y_train_filt (First 20 rows where target == 1) and_
↳top 20k if model predicts perfectly the top 20 churners = {perf}')
plt.xlabel('Index')
plt.ylabel('y_train_filt')
plt.show()

```

Scatter Plot of y_train_filt (First 20 rows where target == 1) and top 20k if model predicts perfectly the top 20 churners = cost 13.176824
dtype: float64



The visualizations and analysis show that focusing on the top 1000 predictions is a practical and effective way to ensure high profits in the top 20.

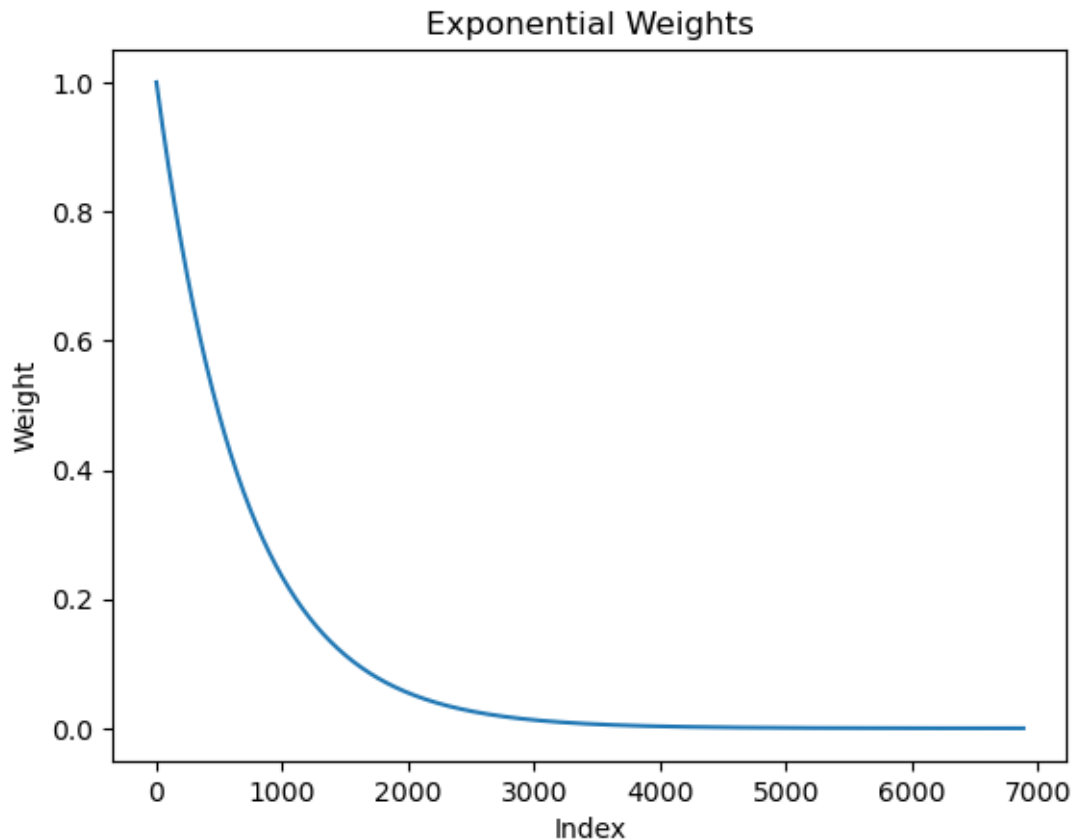
The decision to use a decreasing exponential function for weighting the training data further reinforces this strategy. Exponential weighting places higher importance on the top-ranked data points, aligning the model's focus with the most critical predictions. This method ensures that the model is trained to prioritize and perform exceptionally well on the most valuable subset of predictions, which ultimately drives higher profits in the top 20.

```

[84]: # Generate decreasing exponential weights
weights = np.exp(np.linspace(0, -10, len(y_train_filt)))

```

```
# Plot the weights
plt.plot(weights)
plt.title('Exponential Weights')
plt.xlabel('Index')
plt.ylabel('Weight')
plt.show()
```



Firstly, models set by default parameters are trained for reference.

The following code is designed to evaluate several machine learning models using cross-validation and ROC curves, with an emphasis on adjusting sample weights to emphasize more on top average cost min. It includes models such as Logistic Regression, Random Forest, Bagging Classifier, and XGBoost Classifier. By iterating over different values of k from 1 to 15 every 3 steps, which adjusts the weights applied to the samples using an exponential function, the code systematically explores the impact of varying sample emphasis on model performance.

For each value of k , a new figure is created to plot the ROC curves for each model side-by-side. Within the nested loop, each model is trained and evaluated. Special handling is applied to the XGBoost Classifier due to its requirement for categorical data to be in a specific format.

Cross-validation scores are computed using the ROCmetric, which evaluates the model's ability to distinguish between classes. The models are then fitted with the computed sample weights, and

predictions are made on a validation set.

The evaluation process includes calculating the confusion matrix, which provides insight into the types of errors made by the model, and computing the AUC score, which is averaged over the cross-validation folds. The ROC curve for each model is plotted, displaying the true positive rate against the false positive rate, with the AUC and a custom profit metric shown in the legend. This visualization helps in comparing the models' discriminative abilities.

Overall, this approach provides a comprehensive evaluation of different models under varying conditions of sample weighting, which is crucial for understanding their performance. It ensures that the models are assessed not just on their accuracy, but also on their robustness and ability to handle different emphasis on training samples, thus providing a more reliable measure of their real-world applicability.

```
[87]: models = {
    'Logistic Regression': LogisticRegression(),
    'BaggingClassifier': BaggingClassifier(),
    'XGBClassifier': XGBClassifier(),
}
for k in range(1, 15, 3):
    plt.figure(figsize=(15, 5)) # Create a figure for subplots
    for i, (name, model) in enumerate(models.items(), 1):
        weights = np.exp(np.linspace(0, -k, len(y_train_filt)))
        plt.subplot(1, len(models), i) # Create subplots
        # Use cross-validation to train the model
        if model == XGBClassifier():
            gb = XGBClassifier()
            cats1 = X_train_filt.select_dtypes(exclude=np.number).columns.
→tolist()
            cats2 = X_val.select_dtypes(exclude=np.number).columns.tolist()
            cats3 = X_test.select_dtypes(exclude=np.number).columns.tolist()
            for col in cats1:
                X_train_filt[col] = X_train_filt[col].astype('float')
            for col in cats2:
                X_val[col] = X_val[col].astype('float')
            for col in cats3:
                X_test[col] = X_test[col].astype('float')
            scores = cross_val_score(gb, X_train_filt, y_train_filt, cv=5,
→scoring='roc_auc', fit_params={'sample_weight': weights})
            gb.fit(X_train_filt, y_train_filt, sample_weight=weights)
        else:
            scores = cross_val_score(model, X_train_filt, y_train_filt, cv=5,
→
```



```

→scoring='roc_auc', fit_params={'sample_weight': weights})
    model.fit(X_train_filt, y_train_filt, sample_weight=weights)
    # Predict probabilities for the validation set
    y_pred_proba = model.predict_proba(X_val)
    y_pred = model.predict(X_val)
    pf = profit_top_20(X_val, y_val, y_pred)
    # Compute confusion matrix
    cm = confusion_matrix(y_val, y_pred)
    print(f"Confusion matrix for {name}:")
    print(cm)
    # Calculate AUC
    auc_score = np.mean(scores)
    print(f"AUC for {name}: {auc_score}") # roc_auc
    # Plot ROC curve
    fpr, tpr, _ = roc_curve(y_val, y_pred_proba[:, 1])
    roc_auc = auc(fpr, tpr)
    print('for k in the weights : ', k)
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.
→2f) profit_20 (cost = %0.2f)' % (roc_auc, pf))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC Curve for {name}')
    plt.legend(loc="lower right")
plt.tight_layout() # Adjust the layout to prevent overlapping
plt.show()
'Random Forest': RandomForestClassifier(),

```

Confusion matrix for Logistic Regression:

```
[[787  69]
 [ 33 120]]
```

AUC for Logistic Regression: 0.9173805417066759

for k in the weights : 1

Confusion matrix for Random Forest:

```
[[811  45]
 [ 38 115]]
```

AUC for Random Forest: 0.941425501554083

for k in the weights : 1

Confusion matrix for BaggingClassifier:

```
[[804  52]
 [ 46 107]]
```

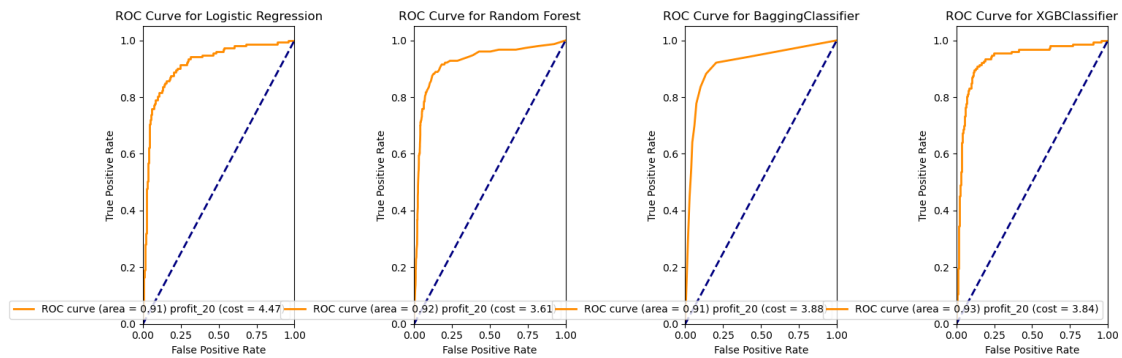
AUC for BaggingClassifier: 0.829544583971795

for k in the weights : 1

Confusion matrix for XGBClassifier:

```
[[805  51]
 [ 40 113]]
```

AUC for XGBClassifier: 0.8979800283427185
for k in the weights : 1



Confusion matrix for Logistic Regression:

```
[[770  86]
 [ 31 122]]
```

AUC for Logistic Regression: 0.9078249363898099

for k in the weights : 4

Confusion matrix for Random Forest:

```
[[816  40]
 [ 40 113]]
```

AUC for Random Forest: 0.9310419089856599

for k in the weights : 4

Confusion matrix for BaggingClassifier:

```
[[812  44]
 [ 40 113]]
```

AUC for BaggingClassifier: 0.7555080669659433

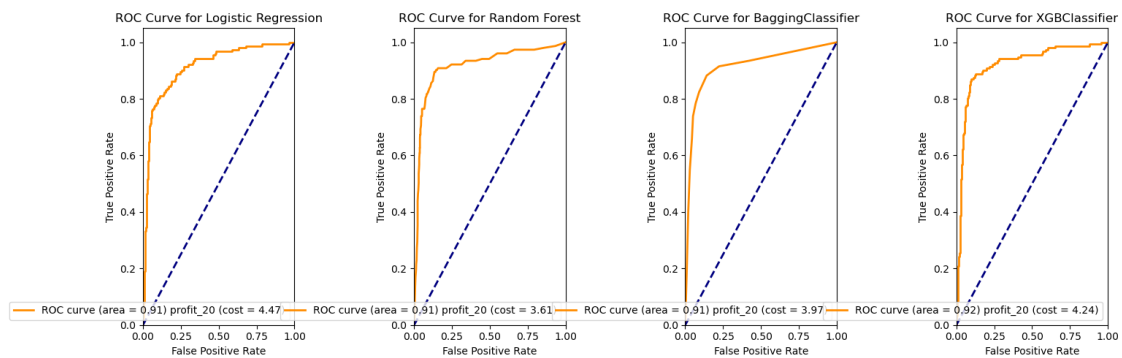
for k in the weights : 4

Confusion matrix for XGBClassifier:

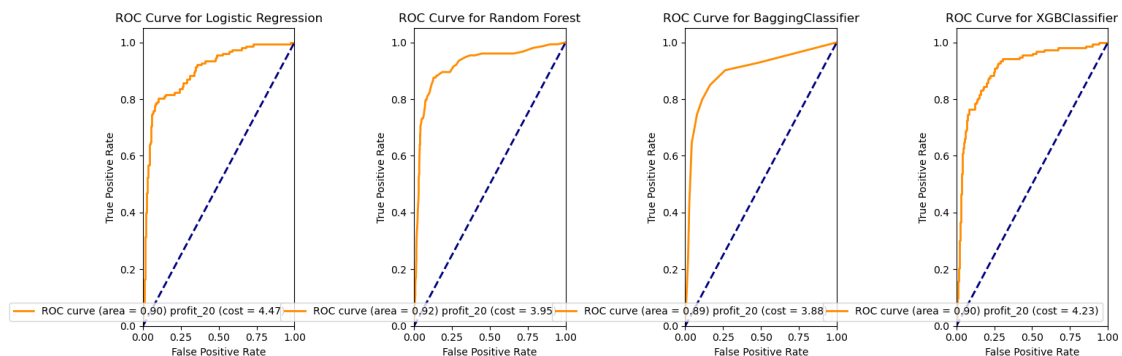
```
[[803  53]
 [ 35 118]]
```

AUC for XGBClassifier: 0.8880053959968557

for k in the weights : 4



Confusion matrix for Logistic Regression:
[[690 166]
[28 125]]
AUC for Logistic Regression: 0.891784084879623
for k in the weights : 7
Confusion matrix for Random Forest:
[[820 36]
[45 108]]
AUC for Random Forest: 0.9274684560130618
for k in the weights : 7
Confusion matrix for BaggingClassifier:
[[807 49]
[48 105]]
AUC for BaggingClassifier: 0.7283890114061121
for k in the weights : 7
Confusion matrix for XGBClassifier:
[[775 81]
[36 117]]
AUC for XGBClassifier: 0.8681374783514213
for k in the weights : 7

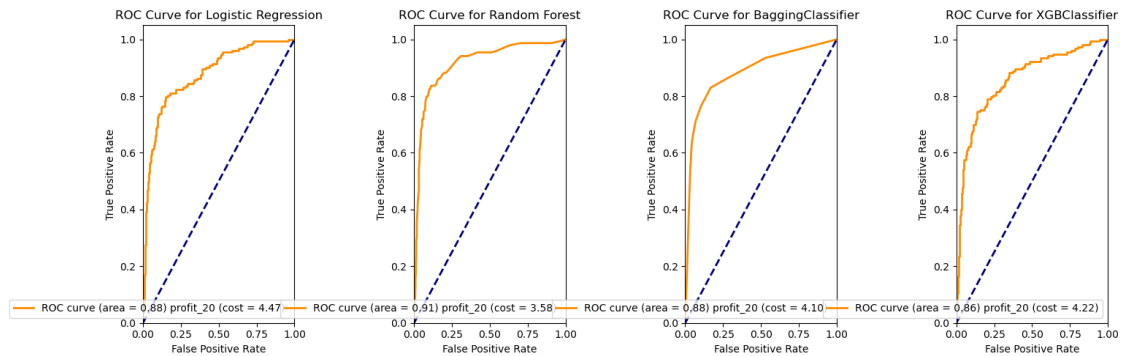


Confusion matrix for Logistic Regression:
[[629 227]
[27 126]]
AUC for Logistic Regression: 0.8735695380047709
for k in the weights : 10
Confusion matrix for Random Forest:
[[814 42]
[48 105]]
AUC for Random Forest: 0.9232027589228299
for k in the weights : 10
Confusion matrix for BaggingClassifier:

```

[[811 45]
 [ 51 102]]
AUC for BaggingClassifier: 0.7325003457967453
for k in the weights : 10
Confusion matrix for XGBClassifier:
[[742 114]
 [ 43 110]]
AUC for XGBClassifier: 0.8067768293327543
for k in the weights : 10

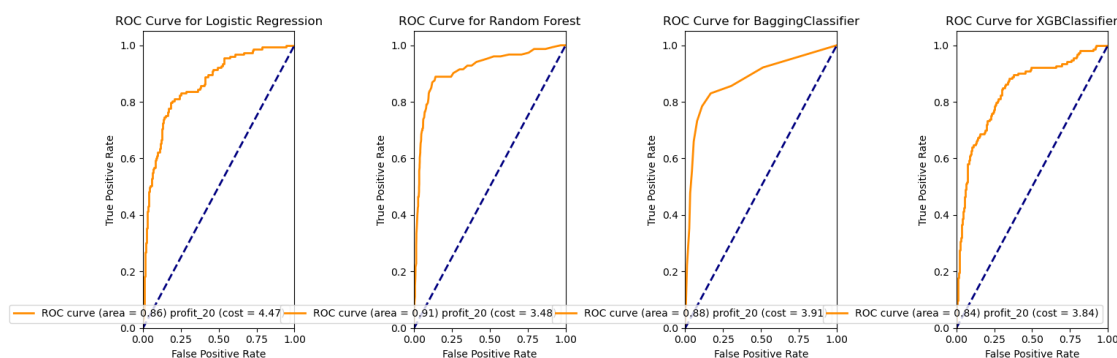
```



```

Confusion matrix for Logistic Regression:
[[561 295]
 [ 25 128]]
AUC for Logistic Regression: 0.8595128035546754
for k in the weights : 13
Confusion matrix for Random Forest:
[[815 41]
 [ 48 105]]
AUC for Random Forest: 0.9253614460841455
for k in the weights : 13
Confusion matrix for BaggingClassifier:
[[808 48]
 [ 52 101]]
AUC for BaggingClassifier: 0.7018853157592736
for k in the weights : 13
Confusion matrix for XGBClassifier:
[[686 170]
 [ 43 110]]
AUC for XGBClassifier: 0.8079590185113871
for k in the weights : 13

```



$$\text{weights} = e^{-10x}$$

[53]: *# final weights*

```
weights = np.exp(np.linspace(0, -10, len(y_train_filt)))
```

After comparing the graphs, we've noticed that increasing the value of `compromises` the AUC and only marginally improves the overall profit top 20. Therefore, we've decided to stick with this final weight as it strikes the best balance between performance and profit. ##### Comparing all the models using their default parameters, it's evident that logistic regression and XGBoost outperform the others in terms of both AUC accuracy and Profit top 20. This is why we chose to further explore only these two models.

0.7 Models interpretations with default parameters

0.7.1 Important features in each model

Important features refer to the specific variables or factors that significantly influence the outcome or prediction of a model. These features are crucial for understanding the underlying patterns and relationships within the data and are often used to make decisions or predictions. In the context of machine learning models like logistic regression and XGBoost, important features are identified based on their impact on the model's performance metrics, such as accuracy, and AUC.

[55]: LR = LogisticRegression()

```
LR.fit(X_train_filt, y_train_filt)
```

```
# Get absolute coefficients for feature importance
```

```
coef_abs = np.abs(LR.coef_[0])
```

```
feature_importances = coef_abs / coef_abs.sum() # Normalize to sum up to 1
```

```

# Plot feature importance for logistic regression model
impo_feat = pd.DataFrame({'feature': X_train_filt.columns, 'importance':
    ↳feature_importances})
impo_feat = impo_feat.sort_values('importance', ascending=False)

# Plot feature importance as bar plot
top_n = 20 # Specify the number of top features to plot
top_n_features = impo_feat['feature'].head(top_n)
top_n_importances = impo_feat['importance'].head(top_n)

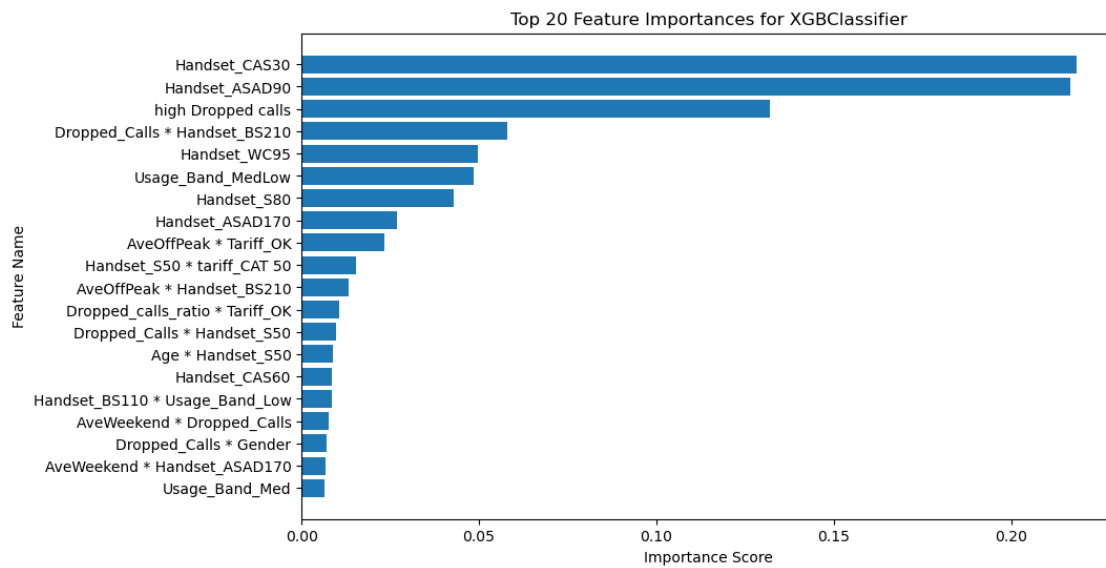
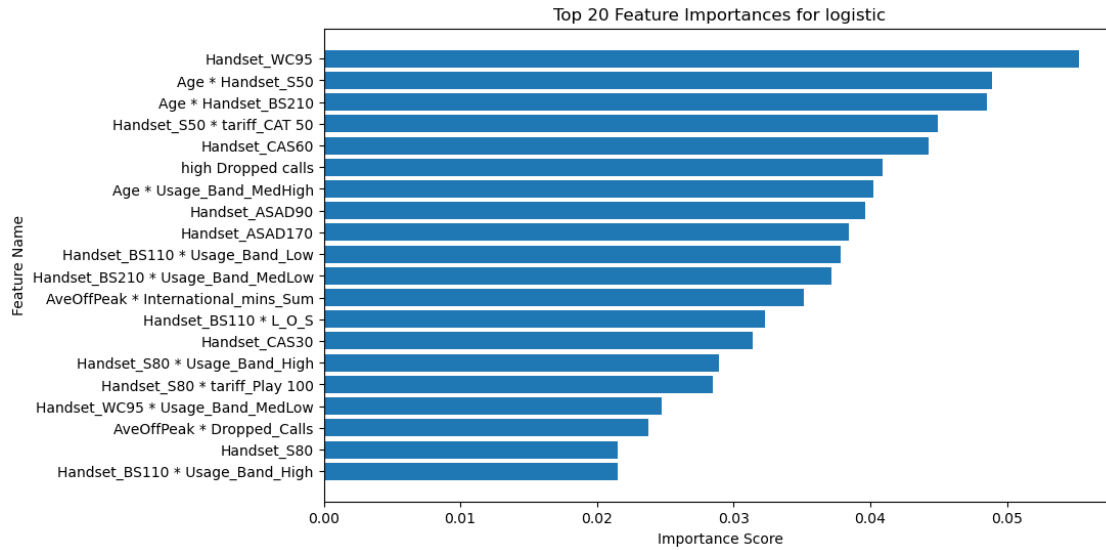
plt.figure(figsize=(10, 6))
plt.barh(top_n_features, top_n_importances)
plt.xlabel('Importance Score')
plt.ylabel('Feature Name')
plt.title(f'Top {top_n} Feature Importances for logistic')
plt.gca().invert_yaxis() # Invert y-axis to display most important features at
    ↳the top
plt.show()

gb = XGBClassifier()
gb.fit(X_train_filt, y_train_filt)
feature_importances = gb.feature_importances_
# Plot feature importance for logistic regression model
impo_feat = pd.DataFrame({'feature': X_train_filt.columns, 'importance':
    ↳feature_importances})
impo_feat = impo_feat.sort_values('importance', ascending=False)

# Plot feature importance as bar plot
top_n = 20 # Specify the number of top features to plot
top_n_features = impo_feat['feature'].head(top_n)
top_n_importances = impo_feat['importance'].head(top_n)

plt.figure(figsize=(10, 6))
plt.barh(top_n_features, top_n_importances)
plt.xlabel('Importance Score')
plt.ylabel('Feature Name')
plt.title(f'Top {top_n} Feature Importances for {name}')
plt.gca().invert_yaxis() # Invert y-axis to display most important features at
    ↳the top
plt.show()

```



The plots display the top 20 most significant features without considering any applied weights. Interestingly, the average minimum cost feature is notably absent from both plots. This observation underscores the importance of feature selection and weighting techniques in model training, as it suggests that the average minimum cost may not have a significant impact on the prediction outcomes compared to other features.

```
[56]: LR = LogisticRegression()
      LR.fit(X_train_filt, y_train_filt, sample_weight= weights )

      coef_abs = np.abs(LR.coef_[0])
```

```

feature_importances = coef_abs / coef_abs.sum()

impo_feat = pd.DataFrame({'feature': X_train_filt.columns, 'importance':
    ↳feature_importances})
impo_feat = impo_feat.sort_values('importance', ascending=False)

top_n = 20
top_n_features = impo_feat['feature'].head(top_n)
top_n_importances = impo_feat['importance'].head(top_n)

plt.figure(figsize=(10, 6))
plt.barh(top_n_features, top_n_importances)
plt.xlabel('Importance Score')
plt.ylabel('Feature Name')
plt.title(f'Top {top_n} Feature Importances for logistic')
plt.gca().invert_yaxis() # Invert y-axis to display most important features at
    ↳the top
plt.show()

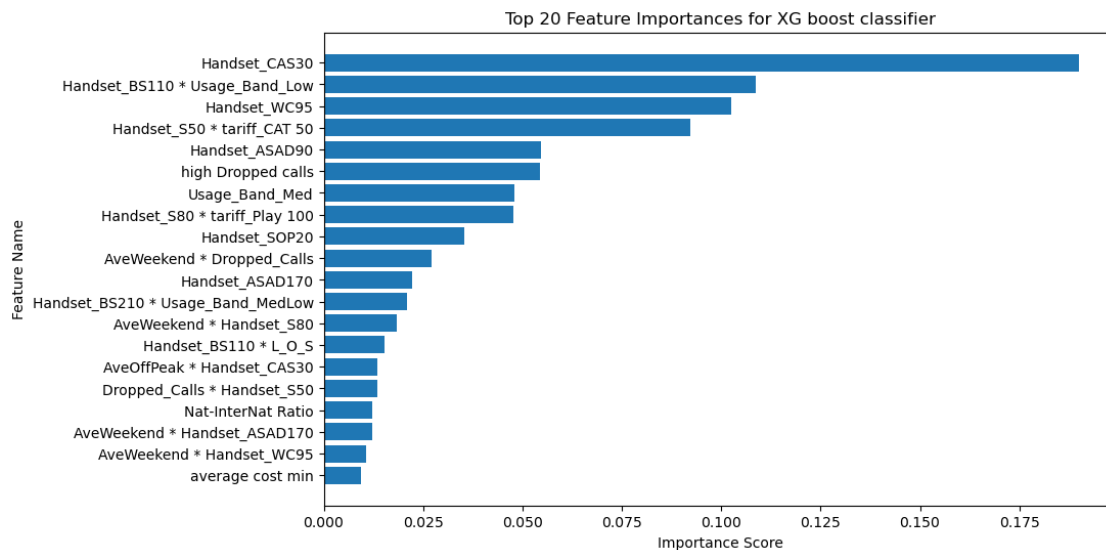
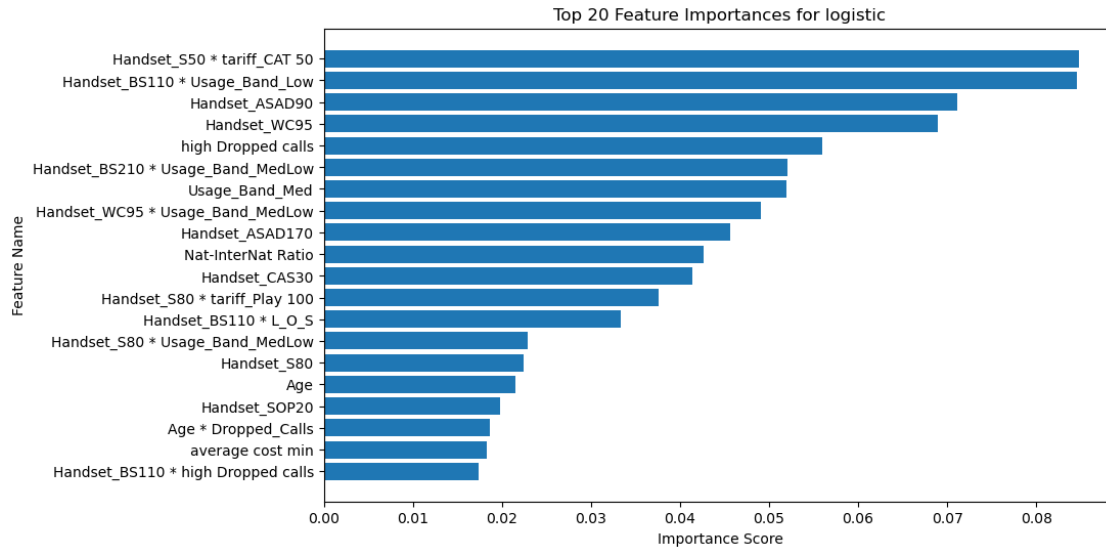
gb.fit(X_train_filt, y_train_filt, sample_weight= weights)
feature_importances = gb.feature_importances_

impo_feat = pd.DataFrame({'feature': X_train_filt.columns, 'importance':
    ↳feature_importances})
impo_feat = impo_feat.sort_values('importance', ascending=False)

top_n = 20
top_n_features = impo_feat['feature'].head(top_n)
top_n_importances = impo_feat['importance'].head(top_n)

plt.figure(figsize=(10, 6))
plt.barh(top_n_features, top_n_importances)
plt.xlabel('Importance Score')
plt.ylabel('Feature Name')
plt.title(f'Top {top_n} Feature Importances for XG boost classifier')
plt.gca().invert_yaxis() # Invert y-axis to display most important features at
    ↳the top
plt.show()

```

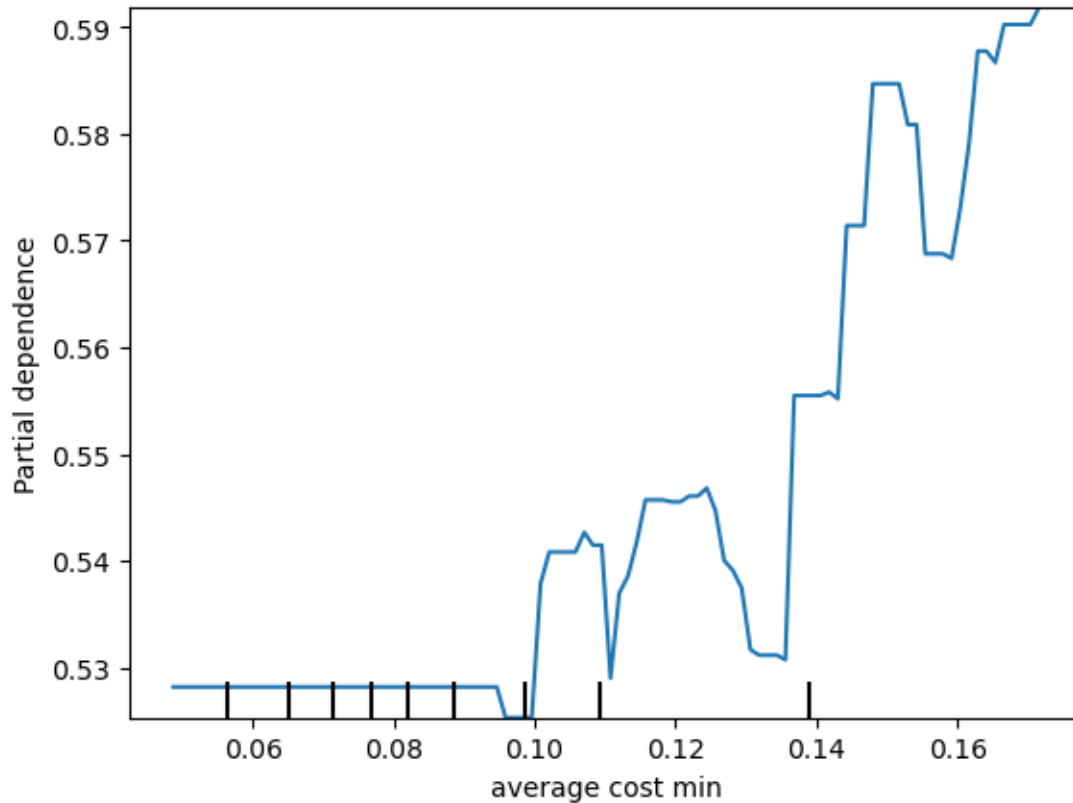



After we added weights, something interesting happened: the “average cost min” feature now shows up in both plots of the top 20 important features. This change suggests that by giving more importance to certain features during model training, we can highlight their significance in making predictions. So, it seems like the “average minimum cost” feature has become more important in helping the model make decisions. This shows how adjusting the weights can improve the model’s performance by focusing on key factors.

This mirrors real-world scenarios where data scientists often prioritize optimizing models based on AUC to highlight overall predictive accuracy. Nevertheless, we also anticipate management’s apprehensions regarding resource limitations.

```
[59]: display = PartialDependenceDisplay.from_estimator(gb, X_train_filt,
↳ features=['average cost min'], kind='average')

plt.show()
```



Partial dependence plots (PDPs) serve as a widely used method for interpreting the correlation between a feature and the predicted outcome within machine learning models. These plots illustrate the marginal impact of a feature on the predicted outcome while considering the collective impacts of all other features.

In our analysis, it's evident that as the average cost min increases, the influence of the 'average cost min' feature on our model becomes more pronounced. By applying exponential weights, we effectively prioritize the instances where the 'average cost min' is greater than 0.1, disregarding those below this threshold.

0.8 Hyperparameter Tuning

First of all, the KFold is designed for randomized search. Due to the imbalance classification, the stratification is adopted. Besides, grid search runs for good especially when computation sources are limited. Instead, randomized search is used for its cost-efficiency. Mostly, randomized search has similar performance with grid search.

A grid search has the advantage that all possible combinations of tuning parameters are considered

and the optimal combination is found. This procedure however becomes extremely time-consuming if a lot of tuning parameters are involved. In such a situation, a randomized search is better to save computation time. A randomized search simply tries m possible combinations out of n cases and returns the best performing one from this subset.

```
[60]: from sklearn.model_selection import RandomizedSearchCV

skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

0.8.1 1. Logistic regression

In logistic regression, most hyperparameters are fixed according to binary classification. The variant hyperparameter could be C . C represents the inverse of regularization strength. Smaller values of C specify stronger regularization. The default value of C is 1 and it must be a positive float.

```
[88]: param_grid = {
        'C': np.arange(0.1, 100.1, 5)
    }

weights = np.exp(np.linspace(0, -10, len(y_train_filt)))

lr = LogisticRegression()
lr_rs = RandomizedSearchCV(lr, param_grid, cv=skf, scoring='roc_auc')
lr_rs.fit(X_train_filt, y_train_filt, sample_weight=weights)

print(lr_rs.best_params_, lr_rs.best_score_)

{'C': 5.1} 0.8778032732371617
```

```
[89]: lr_rs.fit(X_train_filt, y_train_filt, sample_weight=weights)
```

```
[89]: RandomizedSearchCV(cv=StratifiedKFold(n_splits=5, random_state=42,
shuffle=True),
                        estimator=LogisticRegression(),
                        param_distributions={'C': array([ 0.1,  5.1, 10.1, 15.1,
20.1, 25.1, 30.1, 35.1, 40.1, 45.1, 50.1,
55.1, 60.1, 65.1, 70.1, 75.1, 80.1, 85.1, 90.1, 95.1])},
                        scoring='roc_auc')
```

```
[90]: y_pred_Log= lr_rs.predict(X_val) ##### y_pred changes for all

prof_Logisticregression = profit_top_20(X_val, y_pred_Log, y_val)
prof_Logisticregression
```

```
[90]: cost      4.471014
dtype: float64
```

```
[92]: # Lists to store the profit and AUC scores
profit_scores = []
```

```

auc_scores = []

# Iterate over a range of k values
for k in range(1, 100, 3):

    weights = np.exp(np.linspace(0, -k, len(y_train_filt)))

    lr_rs.fit(X_train_filt, y_train_filt, sample_weight=weights)

    y_pred_Log = lr_rs.predict(X_val)
    y_pred_proba_Log = lr_rs.predict_proba(X_val)[: , 1] # Predicted_
    ↪probabilities for AUC

    prof_Logisticregression = profit_top_20(X_val, y_pred_Log, y_val)

    auc_score = roc_auc_score(y_val, y_pred_proba_Log)

    profit_scores.append((k, prof_Logisticregression))
    auc_scores.append((k, auc_score))

    print(f"k={k} | Profit: {prof_Logisticregression} | AUC: {auc_score}")

k_values, profit_values = zip(*profit_scores)
_, auc_values = zip(*auc_scores)

# Plot the profit scores against the k values
plt.figure(figsize=(10, 6))
plt.plot(k_values, profit_values, marker='o', linestyle='-', color='b',
    ↪label='Profit Score')
plt.plot(k_values, auc_values, marker='o', linestyle='-', color='r', label='AUC_
    ↪Score')
plt.xlabel('k')
plt.ylabel('Score')
plt.title('Profit and AUC Scores for Different k Values')
plt.legend()
plt.grid(True)
plt.show()

```

```

k=1 | Profit: cost    4.479056
dtype: float64 | AUC: 0.9130016492578339
k=4 | Profit: cost    4.471014
dtype: float64 | AUC: 0.908512002932014
k=7 | Profit: cost    4.471014
dtype: float64 | AUC: 0.8899807586586036
k=10 | Profit: cost    4.471014

```

```

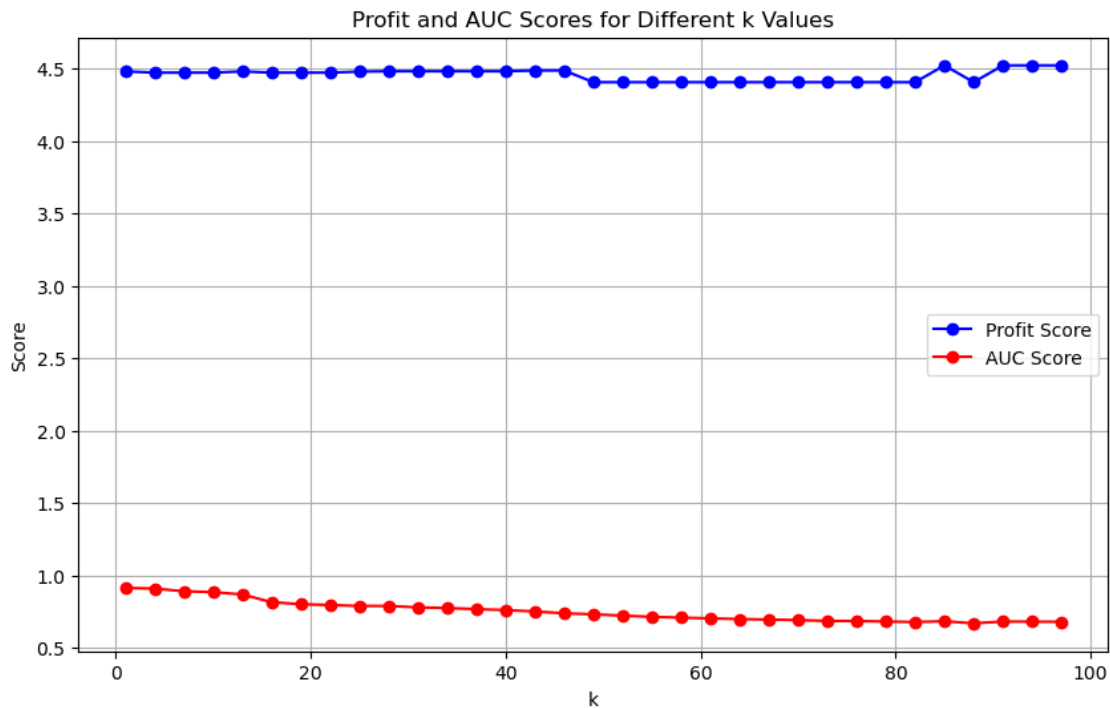
dtype: float64 | AUC: 0.8845824934335104
k=13 | Profit: cost    4.479056
dtype: float64 | AUC: 0.8675477979353734
k=16 | Profit: cost    4.471014
dtype: float64 | AUC: 0.8145806609248061
k=19 | Profit: cost    4.471014
dtype: float64 | AUC: 0.8006459593183066
k=22 | Profit: cost    4.471014
dtype: float64 | AUC: 0.7947666605583044
k=25 | Profit: cost    4.477933
dtype: float64 | AUC: 0.7886124854926395
k=28 | Profit: cost    4.480496
dtype: float64 | AUC: 0.7884903182456782
k=31 | Profit: cost    4.480496
dtype: float64 | AUC: 0.7786863966770509
k=34 | Profit: cost    4.480496
dtype: float64 | AUC: 0.7743418239569971
k=37 | Profit: cost    4.480496
dtype: float64 | AUC: 0.7666376519455134
k=40 | Profit: cost    4.480496
dtype: float64 | AUC: 0.7598268279274327
k=43 | Profit: cost    4.485732
dtype: float64 | AUC: 0.750442856270234
k=46 | Profit: cost    4.485732
dtype: float64 | AUC: 0.7378825361920469
k=49 | Profit: cost    4.405036
dtype: float64 | AUC: 0.7307510231506933
k=52 | Profit: cost    4.405036
dtype: float64 | AUC: 0.7216495632520922
k=55 | Profit: cost    4.405036
dtype: float64 | AUC: 0.7132963777411276
k=58 | Profit: cost    4.405036
dtype: float64 | AUC: 0.7078293934396188
k=61 | Profit: cost    4.405036
dtype: float64 | AUC: 0.7026525563496426
k=64 | Profit: cost    4.405036
dtype: float64 | AUC: 0.6979338464357706
k=67 | Profit: cost    4.405036
dtype: float64 | AUC: 0.6942535581210678
k=70 | Profit: cost    4.405036
dtype: float64 | AUC: 0.691031396982469
k=73 | Profit: cost    4.405036
dtype: float64 | AUC: 0.6851215564107263
k=76 | Profit: cost    4.405036
dtype: float64 | AUC: 0.6848008673874534
k=79 | Profit: cost    4.405036
dtype: float64 | AUC: 0.6806700873495815
k=82 | Profit: cost    4.405036

```

```

dtype: float64 | AUC: 0.6777533443283855
k=85 | Profit: cost    4.521107
dtype: float64 | AUC: 0.6830447132123878
k=88 | Profit: cost    4.405036
dtype: float64 | AUC: 0.6685449880886934
k=91 | Profit: cost    4.521107
dtype: float64 | AUC: 0.6809296927493739
k=94 | Profit: cost    4.521107
dtype: float64 | AUC: 0.6799828965854254
k=97 | Profit: cost    4.521107
dtype: float64 | AUC: 0.6792498931036589

```



0.8.2 2. Boosting - Gradient Boosting

```
[67]: gb = XGBClassifier()
      gb.fit(X_train_filt, y_train_filt, sample_weight=weights)
```

```
[67]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                    colsample_bylevel=None, colsample_bynode=None,
                    colsample_bytree=None, device=None, early_stopping_rounds=None,
                    enable_categorical=False, eval_metric=None, feature_types=None,
                    gamma=None, grow_policy=None, importance_type=None,
                    interaction_constraints=None, learning_rate=None, max_bin=None,
                    max_cat_threshold=None, max_cat_to_onehot=None,
```

```

max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)

```

```

[68]: # hyperparameters tuning
param_grid = {
    'n_estimators': np.arange(1,301,10),
    'max_depth': np.arange(1, 101, 5),
    'learning_rate': np.arange(0.01,1,0.1)
}

# Create RandomizedSearchCV with custom scoring
gb = XGBClassifier()
gb_rs = RandomizedSearchCV(gb, param_grid, cv=skf, scoring='roc_auc', n_jobs=-1)
gb_rs.fit(X_train_filt, y_train_filt, sample_weight=weights)

# Print the best parameters and best score
print(gb_rs.best_params_)

```

```
{'n_estimators': 141, 'max_depth': 31, 'learning_rate': 0.51}
```

```

[69]: # Create a boosting model with the best parameters
gb_rs = XGBClassifier(n_estimators=221, max_depth=21, learning_rate=0.001, gamma=
    ↪ 0.1)

# Fit the model to the training data
gb_rs.fit(X_train_filt, y_train_filt, sample_weight=weights)

```

```

[69]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=0.1, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.001, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=21, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=221, n_jobs=None,
    num_parallel_tree=None, random_state=None, ...)

```

```

[83]: # Initialize lists to store the profit and AUC scores
profit_scores = []
auc_scores = []

# Iterate over a range of k values
for k in range(1, 100, 3):
    # Compute weights

```

```

weights = 100 * np.exp(np.linspace(0, -k, len(y_train_filt)))

# Fit the model
gb_rs.fit(X_train_filt, y_train_filt, sample_weight=weights)

# Predict on the validation set
y_pred_gb_rs = gb_rs.predict(X_val)
y_pred_proba_gb_rs = gb_rs.predict_proba(X_val)[:, 1] # Predicted
↳ probabilities for AUC

# Compute the profit score
prof_gb_rs = profit_top_20(X_val, y_pred_gb_rs, y_val)

# Compute the AUC score
auc_score = roc_auc_score(y_val, y_pred_proba_gb_rs)

# Store the profit score and AUC score along with the corresponding k value
profit_scores.append((k, prof_gb_rs))
auc_scores.append((k, auc_score))

print(f"k={k} | Profit: {prof_gb_rs} | AUC: {auc_score}")

# Separate the k values, profit scores, and AUC scores for plotting
k_values, profit_values = zip(*profit_scores)
_, auc_values = zip(*auc_scores)

# Plot the profit scores and AUC scores against the k values
plt.figure(figsize=(10, 6))
plt.plot(k_values, profit_values, marker='o', linestyle='-', color='b',
↳ label='Profit Score')
plt.plot(k_values, auc_values, marker='o', linestyle='-', color='r', label='AUC
↳ Score')
plt.xlabel('k')
plt.ylabel('Score')
plt.title('Profit and AUC Scores for Different k Values')
plt.legend()
plt.grid(True)
plt.show()

```

```

k=1 | Profit: cost      3.931702
dtype: float64 | AUC: 0.9278220634048012
k=4 | Profit: cost      4.146841
dtype: float64 | AUC: 0.9225688717854743
k=7 | Profit: cost      4.182572
dtype: float64 | AUC: 0.8932869097794881
k=10 | Profit: cost     4.220082
dtype: float64 | AUC: 0.8926531671858775

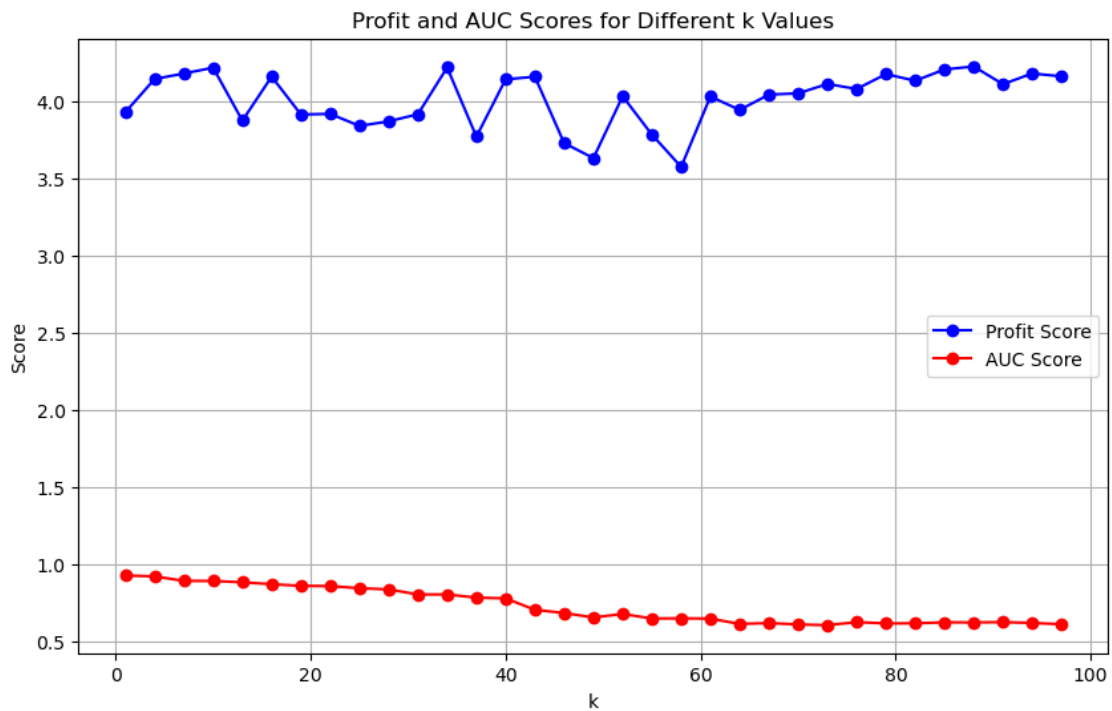
```


k=13 | Profit: cost 3.876909
 dtype: float64 | AUC: 0.884177814427952
 k=16 | Profit: cost 4.162229
 dtype: float64 | AUC: 0.8717702034084662
 k=19 | Profit: cost 3.916937
 dtype: float64 | AUC: 0.8612485492639423
 k=22 | Profit: cost 3.920824
 dtype: float64 | AUC: 0.8593167796713702
 k=25 | Profit: cost 3.84282
 dtype: float64 | AUC: 0.8459012888644554
 k=28 | Profit: cost 3.872283
 dtype: float64 | AUC: 0.8379909596237247
 k=31 | Profit: cost 3.918374
 dtype: float64 | AUC: 0.8051890538146722
 k=34 | Profit: cost 4.221453
 dtype: float64 | AUC: 0.8052806792498931
 k=37 | Profit: cost 3.773306
 dtype: float64 | AUC: 0.7845580599841183
 k=40 | Profit: cost 4.145066
 dtype: float64 | AUC: 0.7797629955408955
 k=43 | Profit: cost 4.161335
 dtype: float64 | AUC: 0.7057372793354101
 k=46 | Profit: cost 3.731266
 dtype: float64 | AUC: 0.6848390446521287
 k=49 | Profit: cost 3.63407
 dtype: float64 | AUC: 0.6561221061633377
 k=52 | Profit: cost 4.03485
 dtype: float64 | AUC: 0.6786848695864638
 k=55 | Profit: cost 3.786174
 dtype: float64 | AUC: 0.6491433021806853
 k=58 | Profit: cost 3.577272
 dtype: float64 | AUC: 0.6504871418972573
 k=61 | Profit: cost 4.033016
 dtype: float64 | AUC: 0.6484484759635942
 k=64 | Profit: cost 3.946624
 dtype: float64 | AUC: 0.6140660313969825
 k=67 | Profit: cost 4.046018
 dtype: float64 | AUC: 0.6196322765866471
 k=70 | Profit: cost 4.05375
 dtype: float64 | AUC: 0.6106835257467473
 k=73 | Profit: cost 4.11482
 dtype: float64 | AUC: 0.6060029930975506
 k=76 | Profit: cost 4.082068
 dtype: float64 | AUC: 0.6257482743876367
 k=79 | Profit: cost 4.17873
 dtype: float64 | AUC: 0.6175325270295033
 k=82 | Profit: cost 4.135503
 dtype: float64 | AUC: 0.6182044468877894

```

k=85 | Profit: cost    4.208585
dtype: float64 | AUC: 0.6243128092358439
k=88 | Profit: cost    4.228098
dtype: float64 | AUC: 0.6240379329301813
k=91 | Profit: cost    4.113219
dtype: float64 | AUC: 0.6257482743876368
k=94 | Profit: cost    4.182785
dtype: float64 | AUC: 0.6205638018447254
k=97 | Profit: cost    4.164162
dtype: float64 | AUC: 0.6119281045751634

```



These plots demonstrate the evolution of AUC and Profit Top 20 as we adjust the parameter “k,” while also considering the impact of our applied weights

0.9 Finale model selection

After conducting a comprehensive comparison between Logistic Regression and XGBoost based on both AUC and Profit Top 20 metrics, we have decided to proceed with XGBoost as our chosen models.

0.10 Reflection

We observed that XGBoost performed worse on the final leaderboard compared to the preliminary one.

To understand why XGBoost performed worse on the final leaderboard compared to the preliminary one, we need to delve into various potential factors that could have contributed to this outcome.

One possibility is that XGBoost may have overfit the training data, performing exceptionally well on the training set but struggling when faced with unseen data. This raises concerns about the model's stability and its ability to generalize to new, unseen instances.

```
[99]: models = {
        "Logistic Regression": lr_rs ,
        "XGBoost": gb_rs
    }

    results = []

    ## with train data

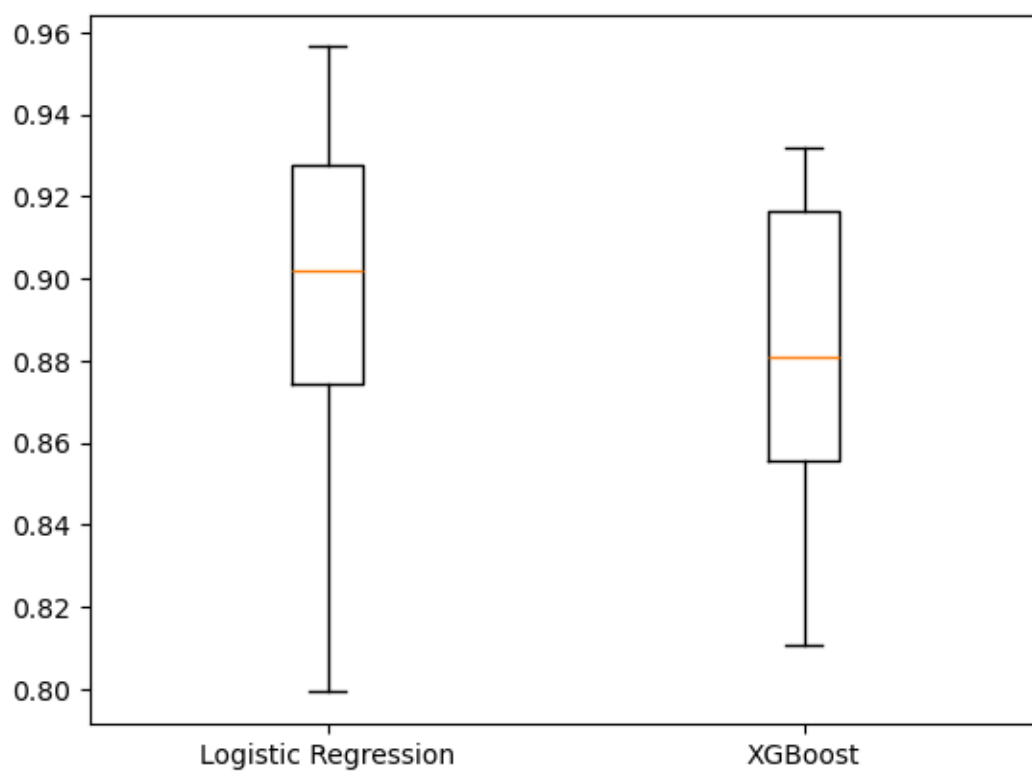
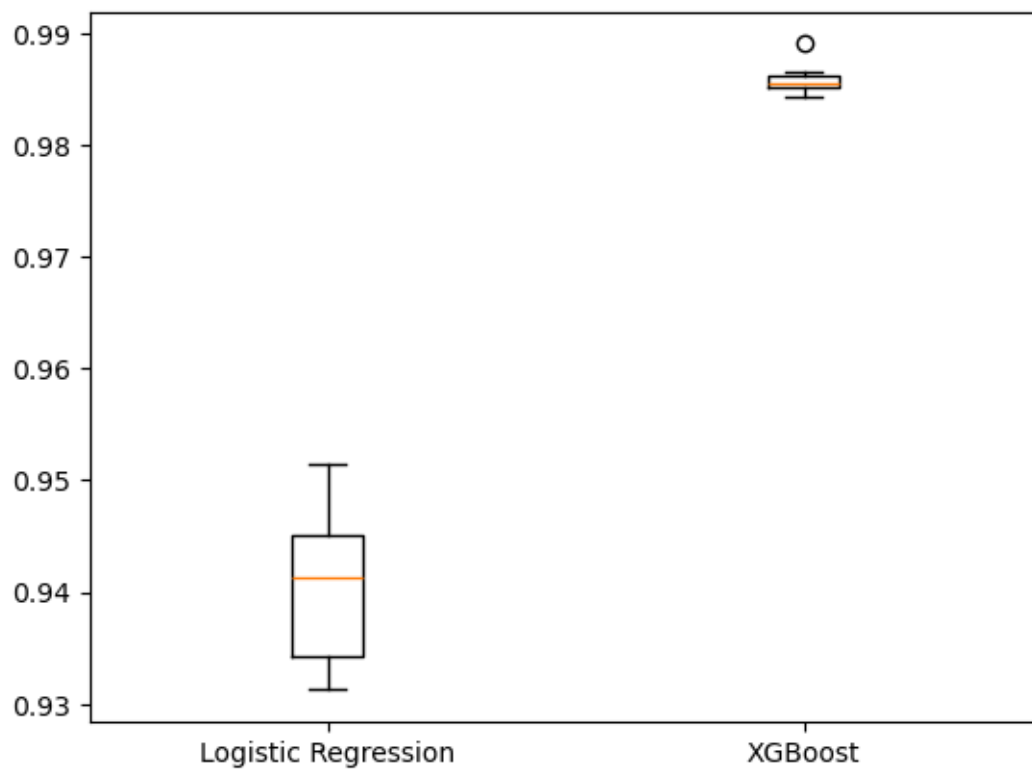
    for model_name, model in models.items():
        kf = KFold(n_splits=6, random_state=42, shuffle=True)
        cv_results = cross_val_score(model, X_train_filt, y_train_filt, cv=kf,
        ↪scoring='roc_auc')
        results.append(cv_results)

    # Plot the results
    plt.boxplot(results, labels=models.keys())
    plt.show()

    #With Test data

    results = []
    for model_name, model in models.items():
        kf = KFold(n_splits=6, random_state=42, shuffle=True)
        cv_results = cross_val_score(model, X_val, y_val, cv=kf, scoring='roc_auc')
        results.append(cv_results)

    # Plot the results
    plt.boxplot(results, labels=models.keys())
    plt.show()
```



To prevent overfitting in XGBoost, the model's performance is penalized by adjusting the parameter gamma, which controls the number of leaf nodes in the tree. The model's objective function combines two parts: the first measures how well the model fits the data, and the second imposes a penalty for the complexity of the tree. Gamma adds a penalty for each additional leaf node, discouraging overly complex trees. By tuning gamma, we adjust the strength of this penalty, ensuring the model remains both accurate and simple, thus avoiding overfitting.

Here, we've added another aspect to our objective function to penalize complexity:

$$\text{Objective} = \text{Loss} + \gamma \times \text{Number of Leaf Nodes}$$

Second possibility is that we initially put a lot of emphasis on predicting the top 20 outcomes accurately. We did this by utilizing weights that decrease rapidly to our data, which helped us prioritize these important predictions. However, we noticed that despite our efforts to maximize profit in the top 20, our overall model performance, measured by AUC, was around 70

This helps explain why our model performed worse on the final leaderboard compared to the preliminary one.

To counter this issue, we made sure to take an AUC higher than 0.89 while aiming a high profit top 20.

In our analysis, we often faced a conflict between making profit top 20 high and getting the most accurate predictions. This happens a lot in real life, where data experts have to make a trade-off between accurate predictions and money/resources. While focusing on getting good predictions can help us understand how well a model works, it might not always help a business make the most money. So, it's important to find a middle ground between making accurate predictions and being practical.

That's why our challenge was to identify the top 20 choices that would bring in the most profit without sacrificing too much on our accuracy measure, AUC. We needed to strike a balance between maximizing profit and maintaining high accuracy in our predictions.

Lastly, in the preprocessing, a few adjustments could have been made. More variables could have been using different feature selection techniques. There were 82 features in the final training set used, and this could have contributed to the overfitting. Furthermore, using correlation as a feature selection technique was not a good idea because we would less likely capture non linear relationships. The connection time column that we've used might've introduced bias since there is censoring, end times are not observed. Also, we could've used normalization to preprocess some numerical features. There were quite a few extreme values but they remained in the data because we thought outliers helped in our prediction. This meant that middle ground observations were less likely to be predicted as a churner.